

H2O.ai

- 2–5 years of experience in AI, machine learning, data science, or software engineering roles.
- Strong understanding of predictive ML (feature engineering, training, evaluation, interpretation).
- Experience building and deploying ML or AI applications in production or production-like environments.
- Exposure to GenAI: LLMs, SLMs, RAG, prompt engineering, or agentic systems.
- Strong Python programming and software engineering fundamentals (testing, debugging, version control, CI/CD).
- Experience building APIs or backend services.
- Basic understanding of cloud platforms and MLOps/LLMOps concepts.
- Good communication and teamwork skills.